

# What the heck is Project Loom?

Deepu K Sasidharan

@deepu105 | [deepu.tech](https://deepu.tech)





# Hi, I'm Deepu K Sasidharan


- *JHipster co-lead developer*
- *Java Champion*
- *Creator of KDash, JDL Studio*
- *Developer Advocate @ Auth0 by Okta*
- *OSS aficionado, polyglot dev, author, speaker*

 @deepu105@mastodon.online

 deepu.tech

 @deepu105

 deepu05



# Concurrency in Java

<https://deepu.tech/concurrency-in-modern-languages-java/>



# JDK Evolution



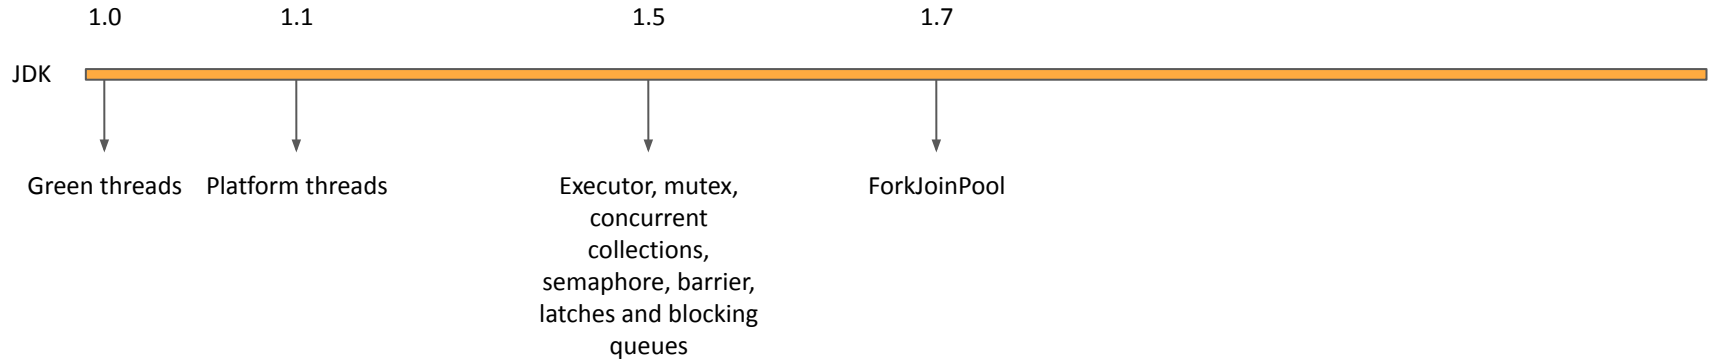


# JDK Evolution



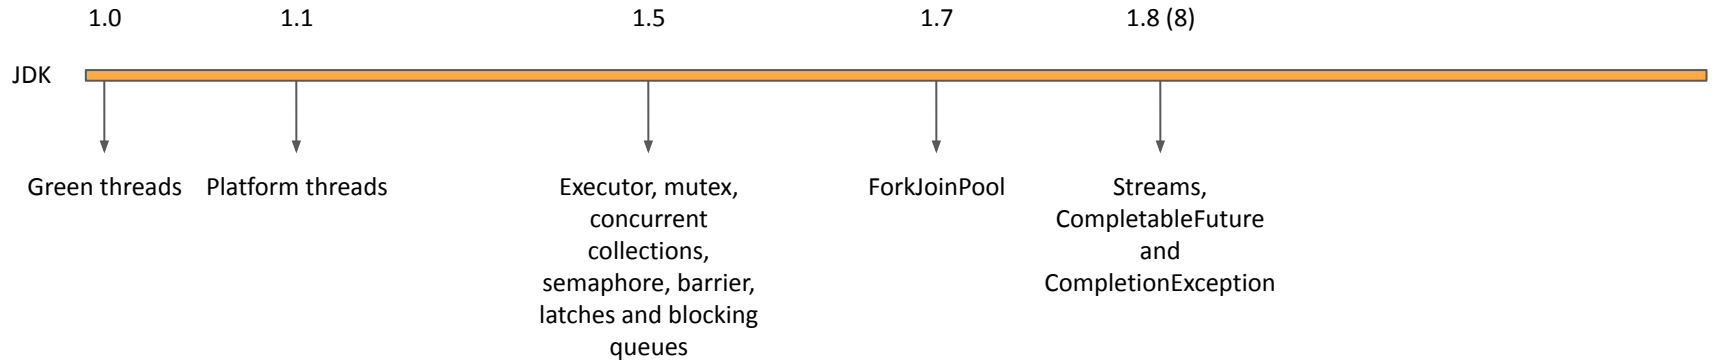


# JDK Evolution



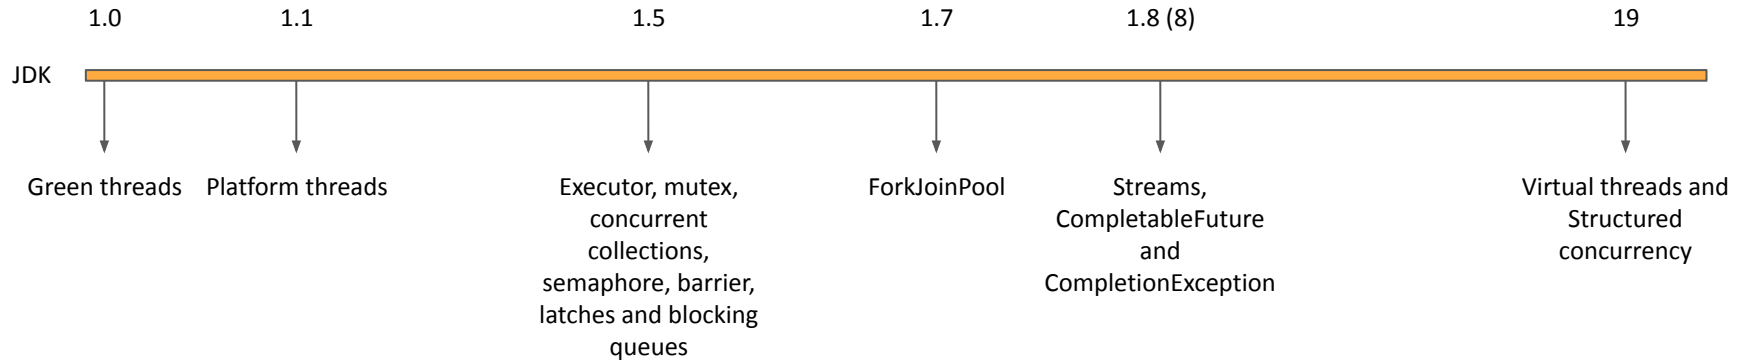


# JDK Evolution





# JDK Evolution





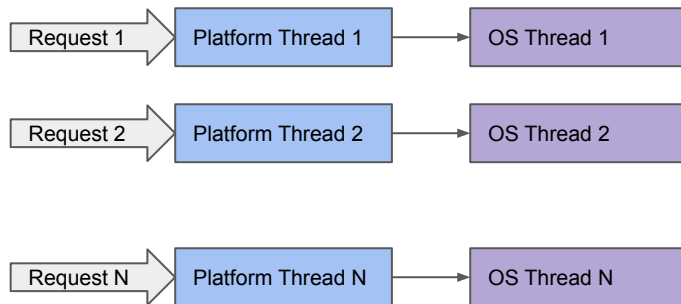


# Platform Threads

- Platform threads == OS threads
- Platforms threads are mapped 1:1 to OS threads




# Thread-per-request model





# Thread-per-request model



## Little's law

$$\lambda = L/W$$


$\lambda$  = Throughput (average rate of requests)

$L$  = Average concurrency (number of requests concurrently processed by the server)

$W$  = Latency (average duration of processing each request)



# Parallel processing



- Should handle data races and data corruption
- Thread synchronization might be needed
- Thread leaks and cancellation delays
- Fragile
- A lot of responsibility on the developer

# Project Loom

<https://developer.okta.com/blog/2022/08/26/state-of-java-project-loom>



# Project Loom

Project Loom aims to drastically reduce the effort of writing, maintaining, and observing high-throughput concurrent applications that make the best use of available hardware.

— Ron Pressler (Tech lead, Project Loom)


# Virtual threads

a.k.a User mode threads

a.k.a Coroutines



# Green threads mapping









# Platform threads mapping

## Green Threads (M:1)




## Platform Threads (1:1)






# Virtual threads mapping


Green Threads (M:1)



Platform Threads (1:1)



Virtual Threads (M:N)





# Goroutines

```
go func() {  
    println("Hello, Goroutines!")  
}()
```



# Kotlin coroutines

```
runBlocking {  
    launch {  
        println("Hello, Kotlin coroutines!")  
    }  
}
```



# Java virtual thread

```
Thread.startVirtualThread(() -> {  
    System.out.println("Hello, Project Loom!");  
});
```



# Virtual thread features

- It is a `Thread` in code, runtime, debugger, and profiler
- It's a Java entity and not a wrapper around an OS thread
- Creating and blocking them are cheap operations
- They should not be pooled
- Virtual threads use a work-stealing `ForkJoinPool` scheduler
- Pluggable schedulers can be used for asynchronous programming
- A virtual thread will have its own stack memory
- The virtual threads API is very similar to platform threads and hence easier to adopt/migrate



# Total number of platform threads

```
var counter = new AtomicInteger();
while (true) {
    new Thread(() -> {
        int count = counter.incrementAndGet();
        System.out.println("Thread count = " + count);
        LockSupport.park();
    }).start();
}
```



# Total number of virtual threads

```
var counter = new AtomicInteger();
while (true) {
    Thread.startVirtualThread(() -> {
        int count = counter.incrementAndGet();
        System.out.println("Thread count = " + count);
        LockSupport.park();
    });
}
```





# Task throughput for platform threads

```
try (var executor = Executors.newThreadPerTaskExecutor(Executors.defaultThreadFactory())) {  
    IntStream.range(0, 100_000).forEach(i -> executor.submit(() -> {  
        Thread.sleep(Duration.ofSeconds(1));  
        System.out.println(i);  
        return i;  
    }));  
}
```

*# 'newThreadPerTaskExecutor' with 'defaultThreadFactory'*

0:18.77 real, 18.15 s user, 7.19 s sys, 135% 3891pu, 0 amem, 743584 mmem

*# 'newCachedThreadPool' with 'defaultThreadFactory'*

0:11.52 real, 13.21 s user, 4.91 s sys, 157% 6019pu, 0 amem, 2215972 mmem



# Task throughput for virtual threads

```
try (var executor = Executors.newVirtualThreadPerTaskExecutor()) {  
    IntStream.range(0, 100_000).forEach(i -> executor.submit(() -> {  
        Thread.sleep(Duration.ofSeconds(1));  
        System.out.println(i);  
        return i;  
    }));  
}
```

0:02.62 real, 6.83 s user, 1.46 s sys, 316% 14840pu, 0 amem, 350268 mmem



# JMH Benchmarks

*# Throughput (more is better)*

Benchmark	Mode	Cnt	Score	Error	Units
LoomBenchmark.platformThreadPerTask	thrpt	5	0.362	± 0.079	ops/s
LoomBenchmark.platformThreadPool	thrpt	5	0.528	± 0.067	ops/s
LoomBenchmark.virtualThreadPerTask	thrpt	5	1.843	± 0.093	ops/s

*# Average time (less is better)*

Benchmark	Mode	Cnt	Score	Error	Units
LoomBenchmark.platformThreadPerTask	avgt	5	5.600	± 0.768	s/op
LoomBenchmark.platformThreadPool	avgt	5	3.887	± 0.717	s/op
LoomBenchmark.virtualThreadPerTask	avgt	5	1.098	± 0.020	s/op

<https://github.com/deepu105/java-loom-benchmarks>



## More benchmarks

- An interesting benchmark using ApacheBench [on GitHub](#) by [Elliot Barlas](#)
- A benchmark using Akka actors [on Medium](#) by [Alexander Zakusylo](#)
- JMH benchmarks for I/O and non-I/O tasks [on GitHub](#) by [Colin Cachia](#)

# Structured concurrency



# Without structured concurrency

```
void handleOrder() throws ExecutionException, InterruptedException {
    try (var esvc = new ScheduledThreadPoolExecutor(8)) {
        Future<Integer> inventory = esvc.submit(() -> updateInventory());
        Future<Integer> order = esvc.submit(() -> updateOrder());

        int theInventory = inventory.get();    // Join updateInventory
        int theOrder = order.get();           // Join updateOrder

        System.out.println("Inventory " + theInventory + " updated for order " + theOrder);
    }
}
```



# Without structured concurrency

```
void handleOrder() throws ExecutionException, InterruptedException {  
    try (var esvc = new ScheduledThreadPoolExecutor(8)) {  
        Future<Integer> inventory = esvc.submit(() -> updateInventory()); // failed  
        Future<Integer> order = esvc.submit(() -> updateOrder()); // runs in background  
  
        int theInventory = inventory.get(); // Join updateInventory // fails  
        int theOrder = order.get(); // Join updateOrder // unreachable  
  
        System.out.println("Inventory " + theInventory + " updated for order " + theOrder);  
    }  
}
```



# Without structured concurrency

```
void handleOrder() throws ExecutionException, InterruptedException {
    try (var esvc = new ScheduledThreadPoolExecutor(8)) {
        Future<Integer> inventory = esvc.submit(() -> updateInventory()); // expensive task
        Future<Integer> order = esvc.submit(() -> updateOrder()); // failed

        int theInventory = inventory.get(); // Join updateInventory // task blocked
        int theOrder = order.get(); // Join updateOrder // will fail

        System.out.println("Inventory " + theInventory + " updated for order " + theOrder);
    }
}
```





# Without structured concurrency

```
void handleOrder() throws ExecutionException, InterruptedException { // interrupted
    try (var esvc = new ScheduledThreadPoolExecutor(8)) {
        Future<Integer> inventory = esvc.submit(() -> updateInventory()); // runs in bg
        Future<Integer> order = esvc.submit(() -> updateOrder()); // runs in bg

        int theInventory = inventory.get(); // Join updateInventory
        int theOrder = order.get(); // Join updateOrder

        System.out.println("Inventory " + theInventory + " updated for order " + theOrder);
    }
}
```



# Structured concurrency

```
void handleOrder() throws ExecutionException, InterruptedException {
    try (var scope = new StructuredTaskScope.ShutdownOnFailure()) {
        Future<Integer> inventory = scope.fork(() -> updateInventory());
        Future<Integer> order = scope.fork(() -> updateOrder());

        scope.join();           // Join both forks
        scope.throwIfFailed();   // ... and propagate errors

        // Here, both forks have succeeded, so compose their results
        System.out.println("Inventory " + inventory.resultNow() + " updated for order " +
            order.resultNow());
    }
}
```



# Structured concurrency

```
void handleOrder() throws ExecutionException, InterruptedException {
    try (var scope = new StructuredTaskScope.ShutdownOnFailure()) {
        Future<Integer> inventory = scope.fork(() -> updateInventory()); // failed
        Future<Integer> order = scope.fork(() -> updateOrder()); // cancelled

        scope.join();           // Join both forks
        scope.throwIfFailed();  // ... and propagate errors

        // Here, both forks have succeeded, so compose their results
        System.out.println("Inventory " + inventory.resultNow() + " updated for order " +
            order.resultNow());
    }
}
```

# State of Project Loom



# Impact for regular developers

- No breaking changes
- Very low API surface and hence easy to adopt/migrate
- Rely on underlying libraries to switch to virtual threads
- Debugging virtual threads would need some getting used to
- Can easily switch to virtual threads from thread pools
- Structured concurrency could help to eliminate a lot of failsafe code
- At the moment need to use preview and incubator modules
- Some unlearning to do (no pooling, no reusing, no shared pool executors)
- Proliferation of virtual threads in simple use cases.



# Impact for libraries

- Performance and throughput increases
- Early adoption
- Simpler codebase
- Server software like tomcat, Undertow and Jetty will see improvements
- Frameworks like Spring, Micronaut and Quarkus will see improvements
- Libraries like RxJava and Akka might benefit from structured concurrency
- Asynchronous and reactive programming will still be around but in many use cases virtual threads could replace them and give same benefits with less complexity



# Early adoption

- GraalVM
  - Support added (<https://github.com/oracle/graal/pull/4802>)
- Quarkus
  - Support added (<https://github.com/quarkusio/quarkus/pull/24942>)
- Micronaut
  - Being discussed  
(<https://github.com/micronaut-projects/micronaut-core/issues/7724>)
- Spring
  - <https://spring.io/blog/2022/10/11/embracing-virtual-threads>

# Caveats





# Resources

- <https://www.infoq.com/articles/java-virtual-threads/>
- <https://inside.java/2020/08/07/loom-performance/>
- [http://cr.openjdk.java.net/~rpressler/loom/loom/sol1\\_part1.html](http://cr.openjdk.java.net/~rpressler/loom/loom/sol1_part1.html)
- <https://foojay.io/today/thinking-about-massive-throughput-meet-virtual-threads/>

# Get the Slides





<https://deepu.tech/tags/java>

# Thank You

Deepu K Sasidharan

@deepu105 | [deepu.tech](https://deepu.tech)



<https://developer.auth0.com>